

Introduction to Data Science

Lecture 3

Manipulating Tabular Data

File Formats

- Considerations for a file format
 - Data model: tabular, hierarchical, array
 - Physical layout
 - Field units and validation
 - Metadata: header, side file, specification, other?
 - Plain text or binary
 - Encoding: ASCII, UTF-8, other?
 - Delimiters and escaping
 - Compression, encryption, checksums?
 - Schema evolution

File Performance

Read/Write time (626 MB tabular file)

	Read Time (Text)	Write Time (Text)	Read Time (Binary)	Write Time (Binary)
Pandas (Python)	36 secs	45 secs	**	**
Scala/Java	18 secs	21 secs	1-6* secs	1-6* secs

Read-Write Times Comparable

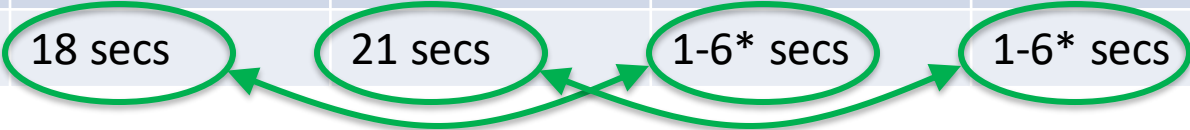
** Pandas doesn't have a default binary file I/O library – you can use Python, but performance depends on what you pick.

* 6 seconds is the time for sustainable read/write.
Often faster due to caching

File Performance

Read/Write time (626 MB tabular file)

	Read Time (Text)	Write Time (Text)	Read Time (Binary)	Write Time (Binary)
Pandas (Python)	36 secs	45 secs	**	**
Scala/Java	18 secs	21 secs	1-6* secs	1-6* secs



Binary I/O much faster than text

** Pandas doesn't have a default binary file I/O library – you can use Python, but performance depends on what you pick.

* 6 seconds is the time for sustainable read/write.
Often faster due to caching

File Performance - Compression

Read/Write time (626 MB tabular file, Scala/Java)

Binary File	Read Time	Write Time	File Size
Gzip level 6 (Java default)	4 secs	75 secs	286 MB
Gzip level 3	4 secs	20 secs	313 MB
Gzip level 1	4 secs	14 secs	328 MB
LZ4 fast	2 secs	4 secs	423 MB
Raw binary file	1-6 secs	1-6 secs	787 MB

Write times much larger than read

Text File	Read Time	Write Time	File Size
Gzip level 6 (default)	26 secs	98 secs	243 MB
Gzip level 3	25 secs	46 secs	259 MB
Gzip level 1	25 secs	33 secs	281 MB
LZ4 fast	22 secs	24 secs	423 MB
Raw text file	18 secs	21 secs	626 MB

File Performance - Compression

Read/Write time (626 MB tabular file, Scala/Java)

Binary File	Read Time	Write Time	File Size
Gzip level 6 (Java default)	4 secs	75 secs	286 MB
Gzip level 3	4 secs	20 secs	313 MB
Gzip level 1	4 secs	14 secs	328 MB
LZ4 fast	2 secs	4 secs	423 MB
Raw binary file	1-6 secs	1-6 secs	787 MB

Large range of compression times



Text File	Read Time	Write Time	File Size
Gzip level 6 (default)	26 secs	98 secs	243 MB
Gzip level 3	25 secs	46 secs	259 MB
Gzip level 1	25 secs	33 secs	281 MB
LZ4 fast	22 secs	24 secs	423 MB
Raw text file	18 secs	21 secs	626 MB

File Performance - Compression

Read/Write time (626 MB tabular file, Scala/Java)

Binary File	Read Time	Write Time	File Size
Gzip level 6 (Java default)	4 secs	75 secs	286 MB
Gzip level 3	4 secs	20 secs	313 MB
Gzip level 1	4 secs	14 secs	328 MB
LZ4 fast	2 secs	4 secs	423 MB
Raw binary file	1-6 secs	1-6 secs	787 MB

Large range of compression times



Text File	Read Time	Write Time	File Size
Gzip level 6 (default)	26 secs	98 secs	243 MB
Gzip level 3	25 secs	46 secs	259 MB
Gzip level 1	25 secs	33 secs	281 MB
LZ4 fast	22 secs	24 secs	423 MB
Raw text file	18 secs	21 secs	626 MB

File Performance - Compression

Read/Write time (626 MB tabular file, Scala/Java)

Binary File	Read Time	Write Time	File Size
Gzip level 6 (Java default)	4 secs	75 secs	286 MB
Gzip level 3	4 secs	20 secs	313 MB
Gzip level 1	4 secs	14 secs	328 MB
LZ4 fast	2 secs	4 secs	423 MB
Raw binary file	1-6 secs	1-6 secs	787 MB

Small range (15%) of
compressed file sizes



Text File	Read Time	Write Time	File Size
Gzip level 6 (default)	26 secs	98 secs	243 MB
Gzip level 3	25 secs	46 secs	259 MB
Gzip level 1	25 secs	33 secs	281 MB
LZ4 fast	22 secs	24 secs	423 MB
Raw text file	18 secs	21 secs	626 MB

File Performance - Compression

Read/Write time (626 MB tabular file, Scala/Java)

Binary File	Read Time	Write Time	File Size
Gzip level 6 (Java default)	4 secs	75 secs	286 MB
Gzip level 3	4 secs	20 secs	313 MB
Gzip level 1	4 secs	14 secs	328 MB
LZ4 fast	2 secs	4 secs	423 MB
Raw binary file	1-6 secs	1-6 secs	787 MB

Text File	Read Time	Write Time	File Size
Gzip level 6 (default)	26 secs	98 secs	243 MB
Gzip level 3	25 secs	46 secs	259 MB
Gzip level 1	25 secs	33 secs	281 MB
LZ4 fast	22 secs	24 secs	423 MB
Raw text file	18 secs	21 secs	626 MB

Binary I/O still much faster than text

File Performance - Compression

Read/Write time (626 MB tabular file, Scala/Java)

Binary File	Read Time	Write Time	File Size
Gzip level 6 (Java default)	4 secs	75 secs	286 MB
Gzip level 3	4 secs	20 secs	313 MB
Gzip level 1	4 secs	14 secs	328 MB
LZ4 fast	2 secs	4 secs	423 MB
Raw binary file	1-6 secs	1-6 secs	787 MB

Text File	Read Time	Write Time	File Size
Gzip level 6 (default)	26 secs	98 secs	243 MB
Gzip level 3	25 secs	46 secs	259 MB
Gzip level 1	25 secs	33 secs	281 MB
LZ4 fast	22 secs	24 secs	423 MB
Raw text file	18 secs	21 secs	626 MB

Binary I/O still much faster than text

File Performance - Compression

Read/Write time (626 MB tabular file, Scala/Java)

Binary File	Read Time	Write Time	File Size
Gzip level 6 (Java default)	4 secs	75 secs	286 MB
Gzip level 3	4 secs	20 secs	313 MB
Gzip level 1	4 secs	14 secs	328 MB
LZ4 fast	2 secs	4 secs	423 MB
Raw binary file	1-6 secs	1-6 secs	787 MB

LZ4 compression \approx raw I/O speed

Text File	Read Time	Write Time	File Size
Gzip level 6 (default)	26 secs	98 secs	243 MB
Gzip level 3	25 secs	46 secs	259 MB
Gzip level 1	25 secs	33 secs	281 MB
LZ4 fast	22 secs	24 secs	423 MB
Raw text file	18 secs	21 secs	626 MB

File Performance - Compression

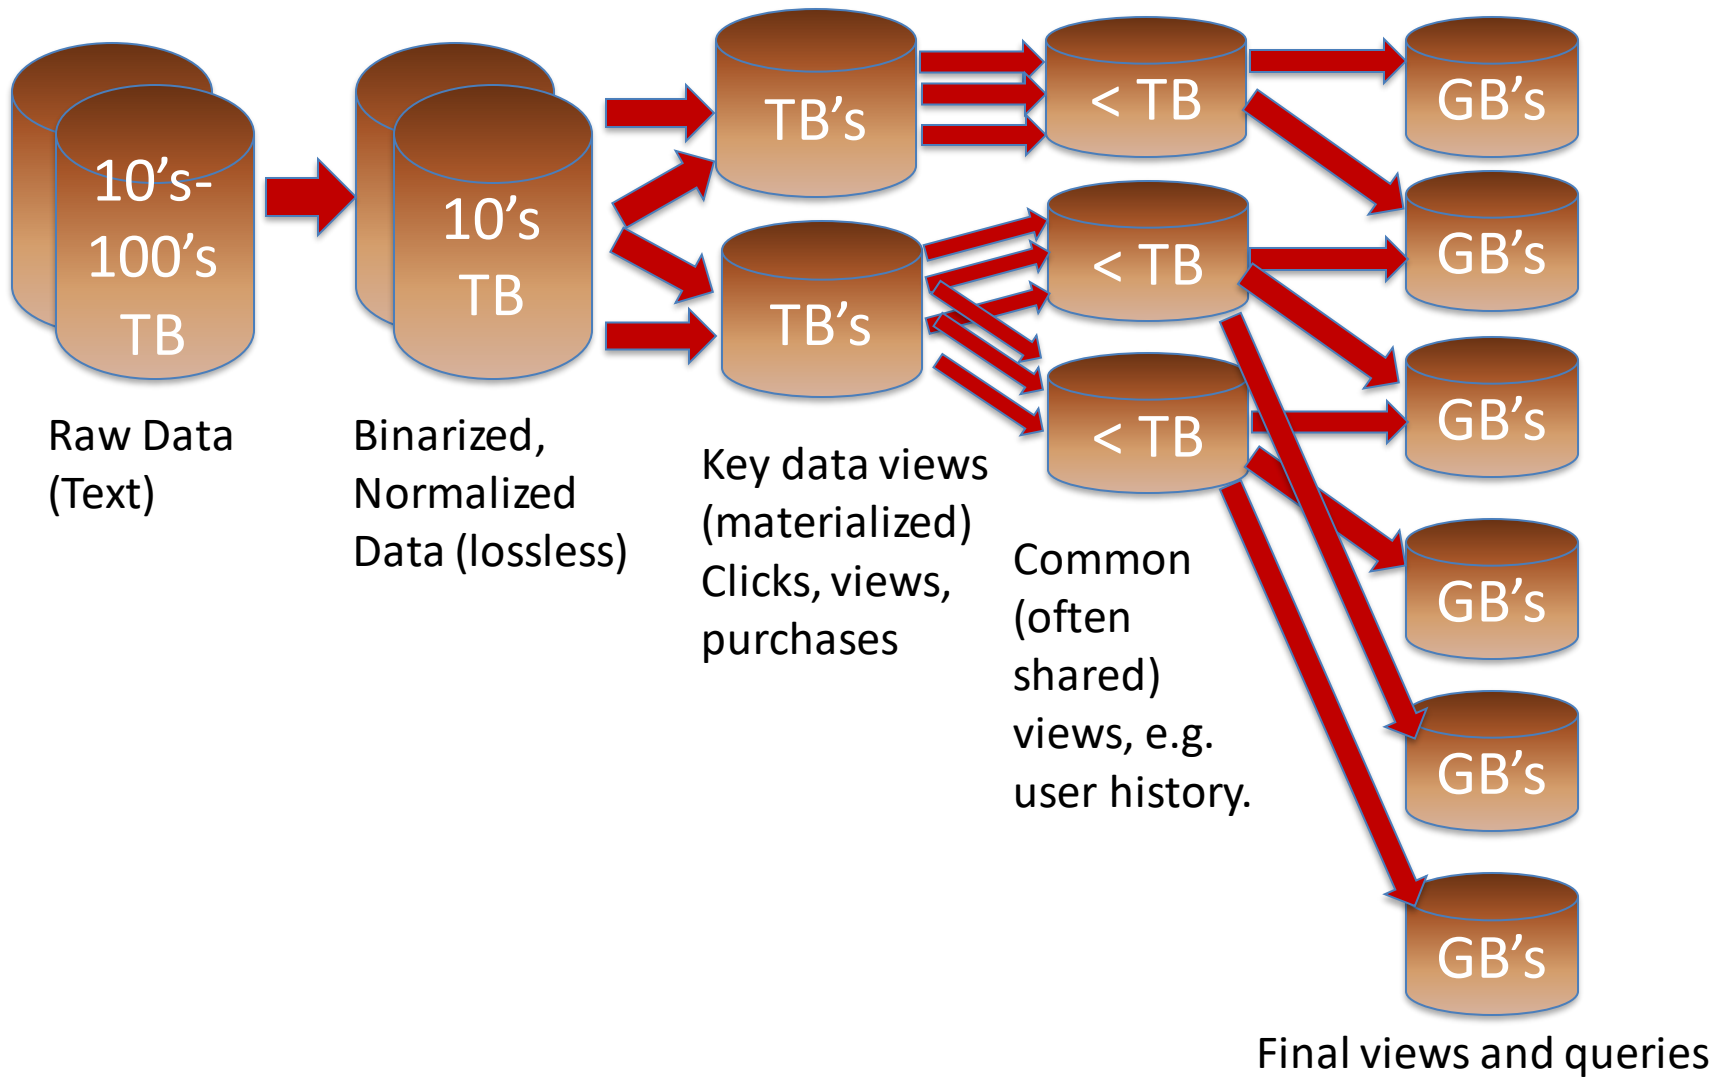
Read/Write time (626 MB tabular file, Scala/Java)

Binary File	Read Time	Write Time	File Size
Gzip level 6 (Java default)	4 secs	75 secs	286 MB
Gzip level 3	4 secs	20 secs	313 MB
Gzip level 1	4 secs	14 secs	328 MB
LZ4 fast	2 secs	4 secs	423 MB
Raw binary file	1-6 secs	1-6 secs	787 MB

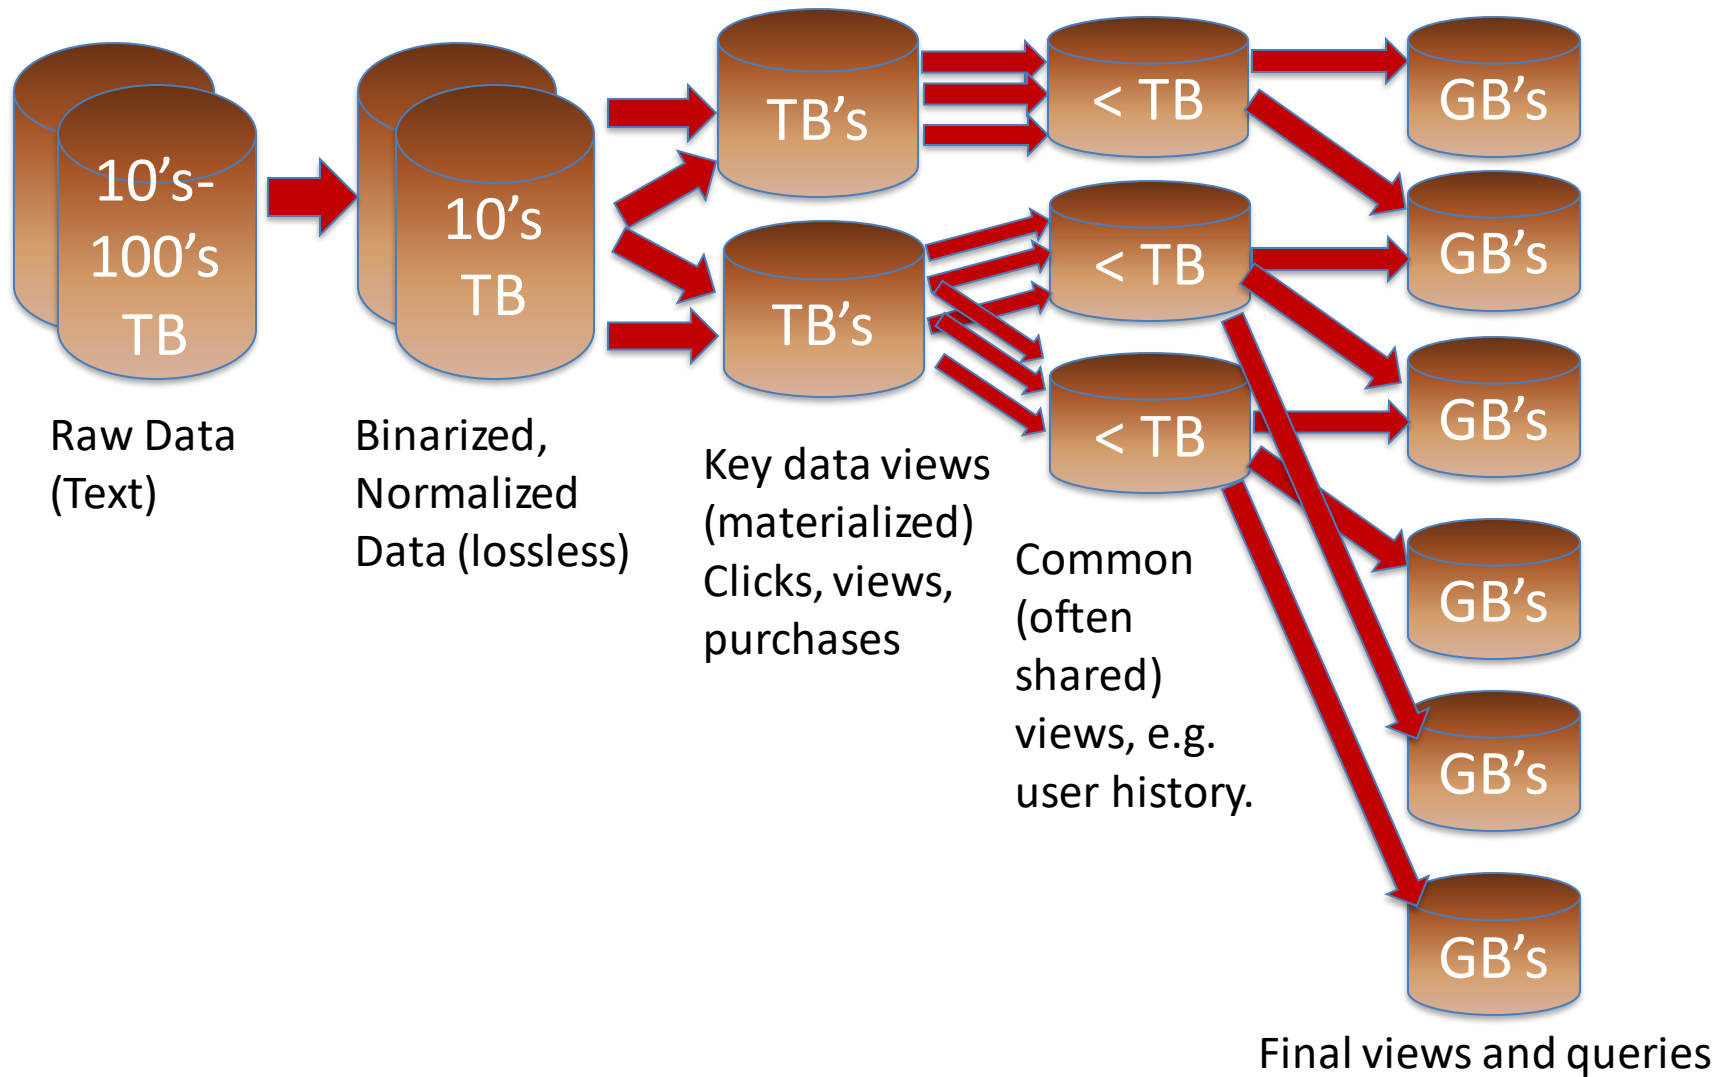
LZ4 compression \approx raw I/O speed

Text File	Read Time	Write Time	File Size
Gzip level 6 (default)	26 secs	98 secs	243 MB
Gzip level 3	25 secs	46 secs	259 MB
Gzip level 1	25 secs	33 secs	281 MB
LZ4 fast	22 secs	24 secs	423 MB
Raw text file	18 secs	21 secs	626 MB


Data Pipeline Design



Data Pipeline Design



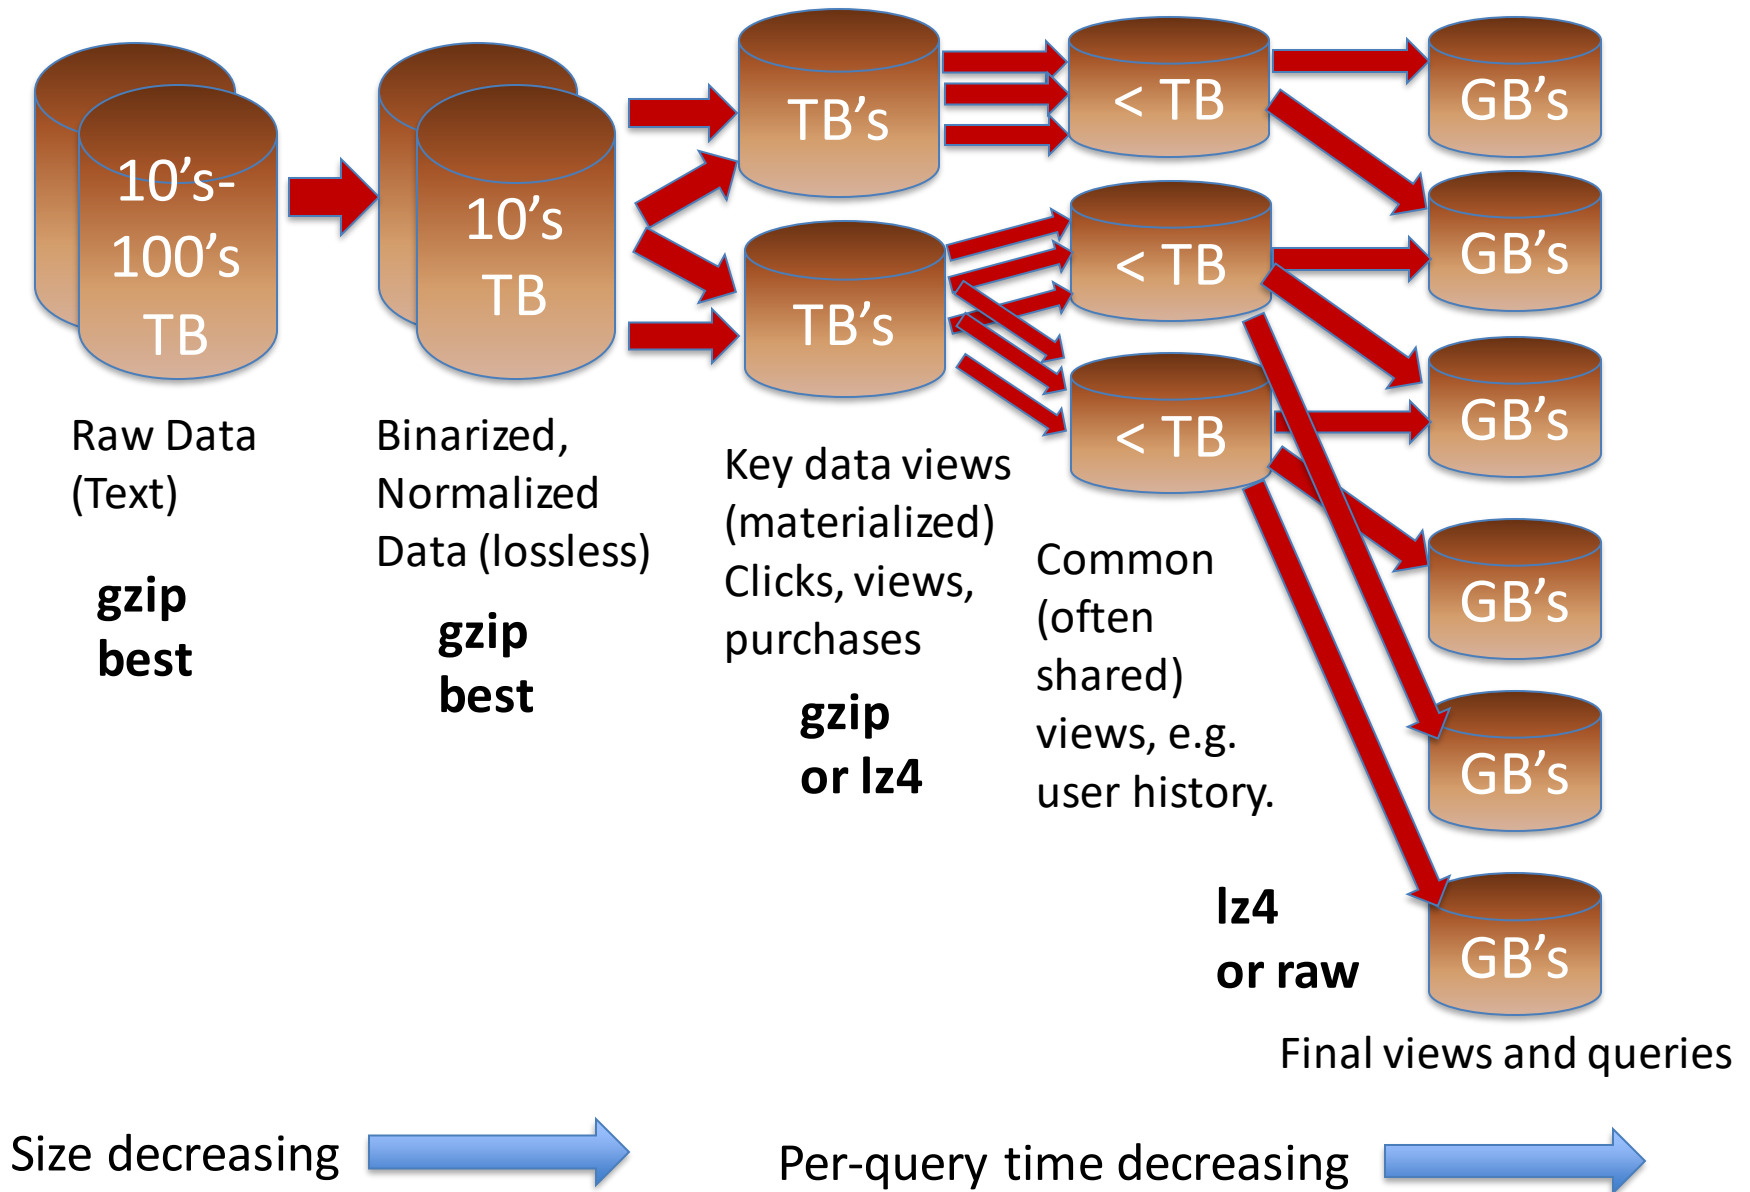
Number of queries Increasing



Updates Increasing



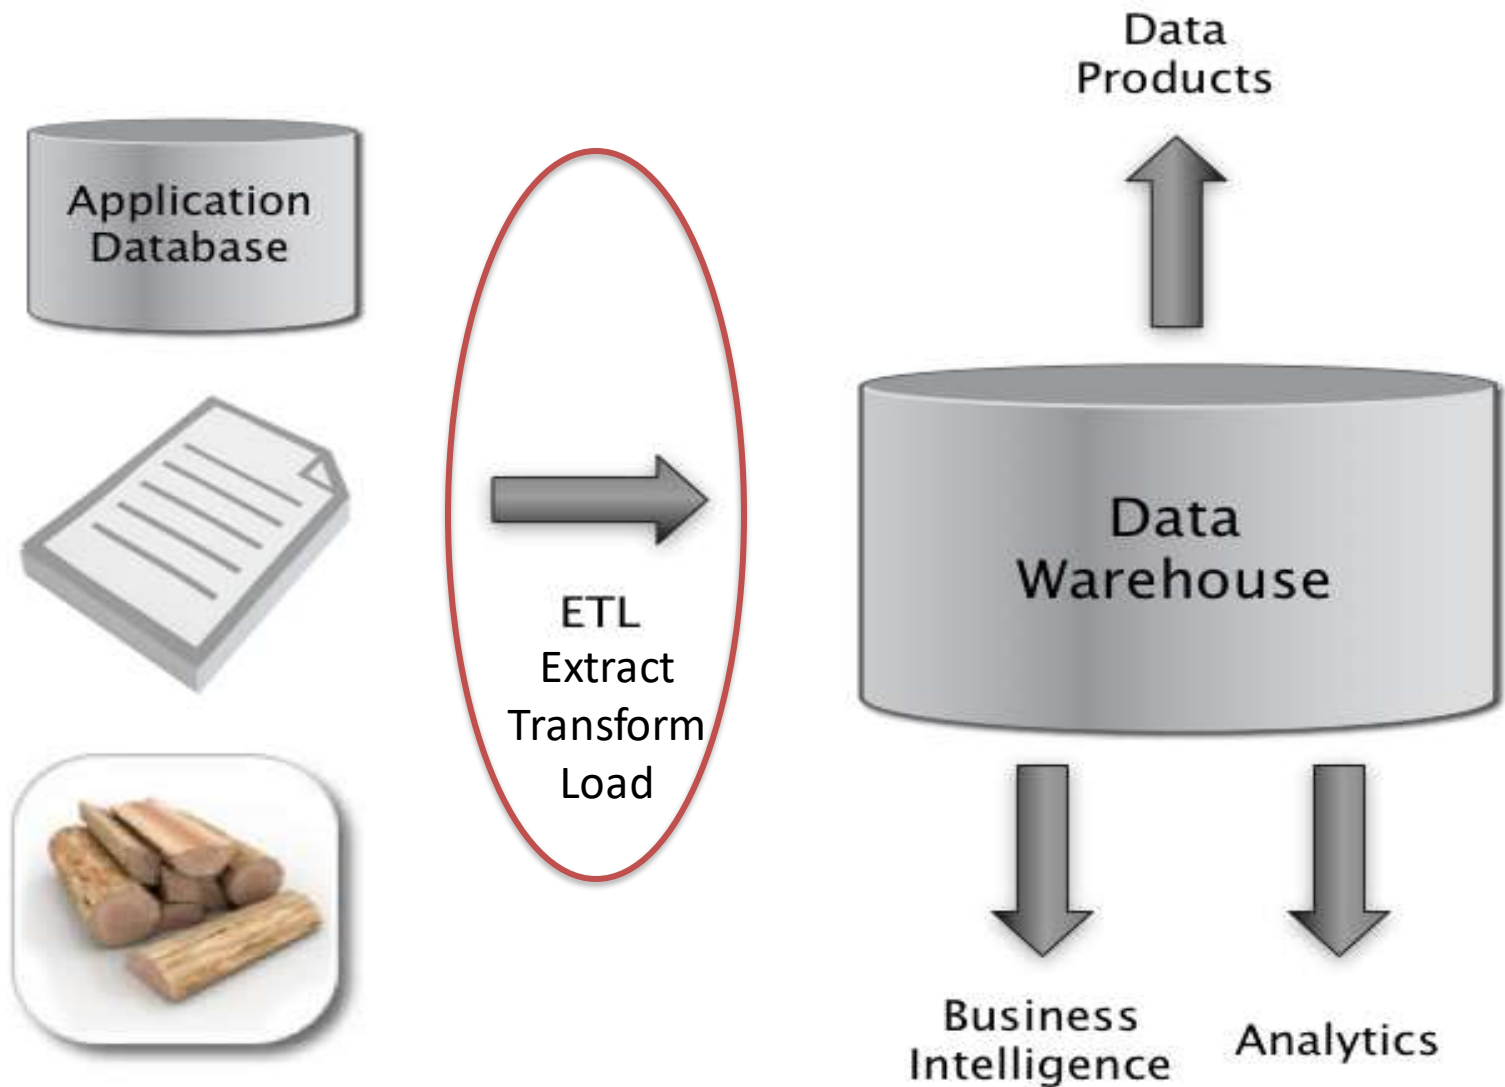
Data Pipeline Design



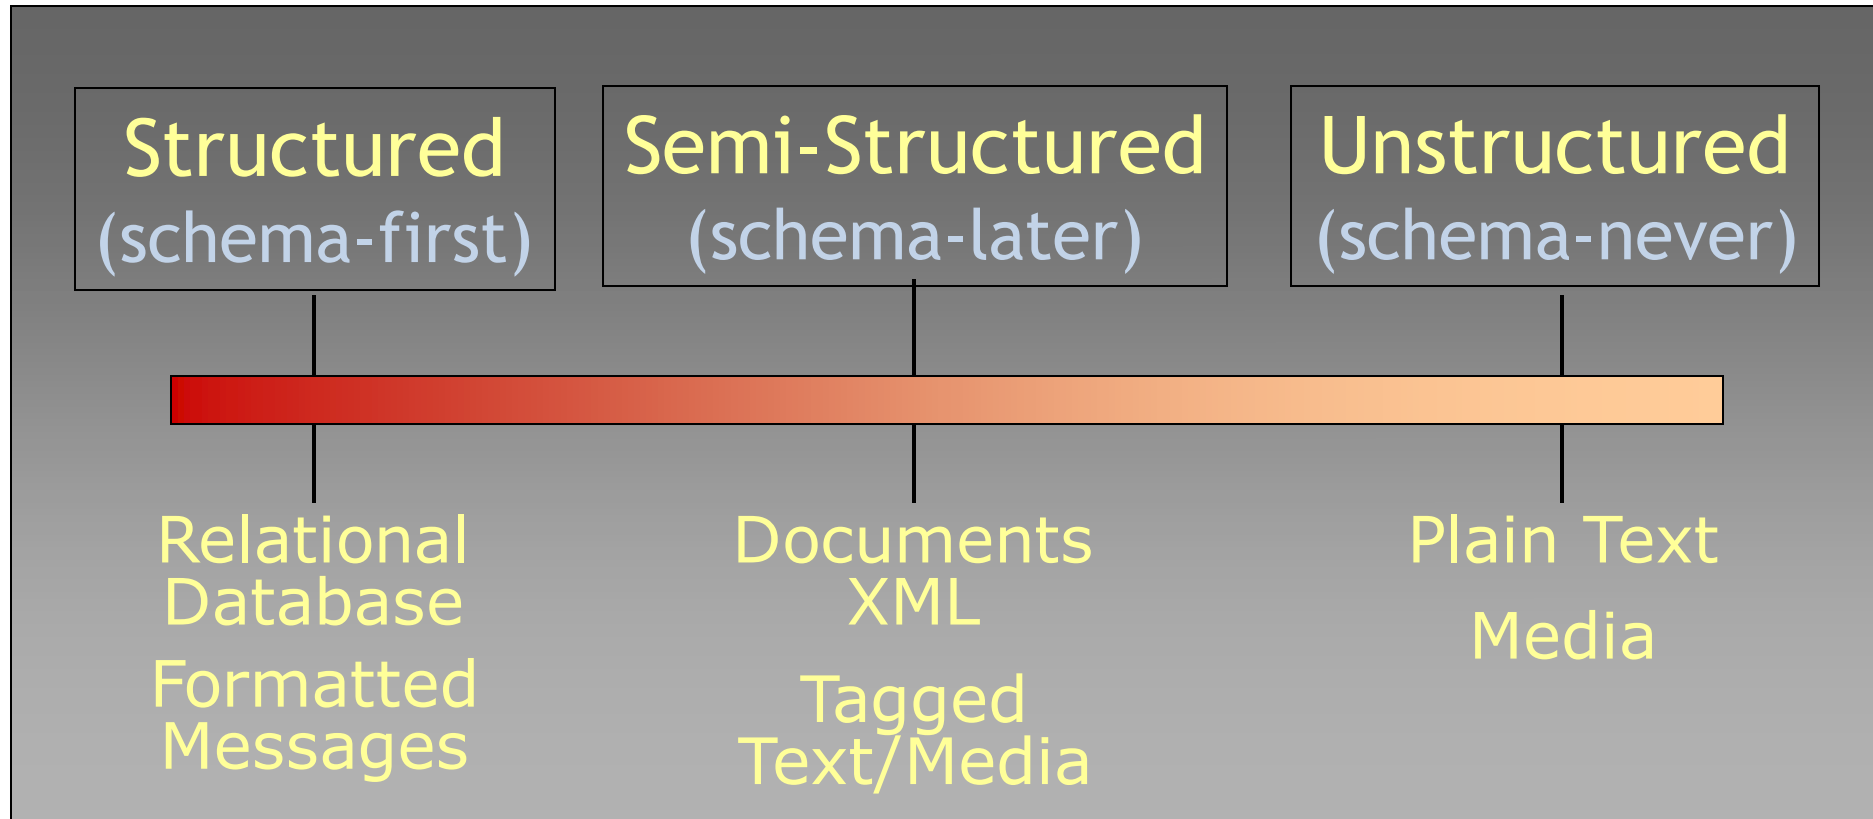
Outline for this Evening

- Data Models, Tables, Structure, etc.
 - SQL
 - NoSQL
 - Schema on Read vs. Schema on Write

The Big Picture



The Structure Spectrum




Key Concept: Structured Data

A *data model* is a collection of concepts for describing data.

A *schema* is a description of a particular collection of data, using a given data model.

The Relational Model*

- The Relational Model is Ubiquitous:
 - MySQL, PostgreSQL, Oracle, DB2, SQLServer, ...
 - Foundational work done at
 - IBM - System R
 - UC Berkeley - Ingres
- Object-oriented concepts have been merged in
 - Early work: POSTGRES research project at Berkeley
 - Informix, IBM DB2, Oracle 8i
- Also has support for XML (semi-structured data)



E. F., "Ted" Codd
Turing Award 1981

**Codd, E. F. (1970). "A relational model of data for large shared data banks".
Communications of the ACM 13 (6): 37*

Relational Database: Definitions

- *Relational database*: a set of *relations*
- *Relation*: made up of 2 parts:

Schema : specifies name of relation, plus name and type of each column

Students(*sid*: string, *name*: string, *login*: string, *age*: integer, *gpa*: real)

Instance : the actual data at a given time

- #rows = *cardinality*
- #fields = *degree / arity*
- A relation is a mathematical object (from set theory) which is true for certain arguments.
- An instance defines the set of arguments for which the relation is true.

Ex: Instance of Students Relation

sid	name	login	age	gpa
53666	Jones	jones@cs	18	3.4
53688	Smith	smith@eecs	18	3.2
53650	Smith	smith @math	19	3.8

- Cardinality = 3, arity = 5 , all rows distinct
- The relation is true for these tuples and false for others

SQL - A language for Relational DBs*

- SQL = Structured Query Language
- Data Definition Language (DDL)
 - create, modify, delete relations
 - specify constraints
 - administer users, security, etc.
- Data Manipulation Language (DML)
 - Specify queries to find tuples that satisfy criteria
 - add, modify, remove tuples
- The DBMS is responsible for efficient evaluation.

* Developed at [IBM](#) by [Donald D. Chamberlin](#) and [Raymond F. Boyce](#) in the 1970s.
Used to be *SEQUEL* (*Structured English QUery Language*)

Creating Relations in SQL

- Create the Students relation.
 - Note: the type (domain) of each field is specified, and enforced by the DBMS whenever tuples are added or modified.

```
CREATE TABLE Students
(sid CHAR(20),
 name CHAR(20),
 login CHAR(10),
 age INTEGER,
 gpa FLOAT)
```


Table Creation (continued)

- Another example: the Enrolled table holds information about courses students take.

```
CREATE TABLE Enrolled  
  (sid CHAR(20),  
   cid CHAR(20),  
   grade CHAR(2))
```

Adding and Deleting Tuples

- Can insert a single tuple using:

```
INSERT INTO Students (sid, name, login, age, gpa)
VALUES ('53688', 'Smith', 'smith@ee', 18, 3.2)
```

- Can delete all tuples satisfying some condition (e.g., name = Smith):

```
DELETE
FROM Students S
WHERE S.name = 'Smith'
```

Queries in SQL

- Single-table queries are straightforward.
- To find all 18 year old students, we can write:

```
SELECT *  
FROM Students S  
WHERE S.age=18
```

- To find just names and logins, replace the first line:

```
SELECT S.name, S.login
```

Querying Multiple Relations

- Can specify a join over two tables as follows:

```
SELECT S.name, E.cid  
FROM Students S, Enrolled E  
WHERE S.sid=E.sid AND E.grade='B'
```

S

sid	cid	grade
53831	Carnatic101	C
53831	Reggae203	B
53650	Topology112	A
53666	History105	B

E

sid	name	login	age	gpa
53831	Jones	jones@cs	18	3.4
53831	Smith	smith@ee	18	3.2

result =

S.name	E.cid
Jones	History105

Note: no referential integrity constraints have been used here.

Basic SQL Query

SELECT	[<i>DISTINCT</i>]	<i>target-list</i>
FROM		<i>relation-list</i>
WHERE		<i>qualification</i>

- *relation-list* : A list of relation names
 - possibly with a *range-variable* after each name
- *target-list* : A list of attributes of tables in *relation-list*
- *qualification* : Comparisons combined using AND, OR and NOT.
 - Comparisons are $\text{Attr } op \text{ const}$ or $\text{Attr1 } op \text{ Attr2}$, where *op* is one of $= \neq < > \leq \geq$
- *DISTINCT*: optional keyword indicating that the answer should not contain duplicates.
 - In SQL SELECT, the default is that duplicates are not eliminated! (Result is called a “multiset”)

SQL Inner Joins

```
SELECT S.name, E.classid  
FROM students S (INNER) JOIN Enrolled E  
ON S.sid=E.sid
```

S

S.name	S.sid
Jones	11111
Smith	22222
Brown	33333

S.name	E.classid
Jones	History105
Jones	DataScience194
Smith	French150

E

E.sid	E.classid
11111	History105
11111	DataScience194
22222	French150
44444	English10

Note the previous version of this query (with no join keyword) is an “Implicit join”

SQL Inner Joins

```
SELECT S.name, E.classid  
FROM Students S (INNER) JOIN Enrolled E  
ON S.sid=E.sid
```

S

S.name	S.sid
Jones	11111
Smith	22222
Brown	33333

S.name	E.classid
Jones	History105
Jones	DataScience194
Smith	French150

E

E.sid	E.classid
11111	History105
11111	DataScience194
22222	French150
44444	English10

Unmatched keys

What kind of Join is this?

```
SELECT S.name, E.classid
FROM Students S ?? Enrolled E
ON S.sid=E.sid
```

S

S.name	S.sid
Jones	11111
Smith	22222
Brown	33333

S.name	E.classid
Jones	History105
Jones	DataScience194
Smith	French150
Brown	NULL

E

E.sid	E.classid
11111	History105
11111	DataScience194
22222	French150
44444	English10

SQL Joins

```
SELECT S.name, E.classid
FROM Students S LEFT OUTER JOIN Enrolled E
ON S.sid=E.sid
```

S

S.name	S.sid
Jones	11111
Smith	22222
Brown	33333

S.name	E.classid
Jones	History105
Jones	DataScience194
Smith	French150
Brown	NULL

E

E.sid	E.classid
11111	History105
11111	DataScience194
22222	French150
44444	English10

What kind of Join is this?

```
SELECT S.name, E.classid
FROM Students S ?? Enrolled E
ON S.sid=E.sid
```

S

S.name	S.sid
Jones	11111
Smith	22222
Brown	33333

S.name	E.classid
Jones	History105
Jones	DataScience194
Smith	French150
NULL	English10

E

E.sid	E.classid
11111	History105
11111	DataScience194
22222	French150
44444	English10

SQL Joins

```
SELECT S.name, E.classid
FROM Students S RIGHT OUTER JOIN Enrolled E
ON S.sid=E.sid
```

S	S.name	S.sid
	Jones	11111
	Smith	22222
	Brown	33333

S.name	E.classid
Jones	History105
Jones	DataScience194
Smith	French150
NULL	English10

E	E.sid	E.classid
	11111	History105
	11111	DataScience194
	22222	French150
	44444	English10

What kind of Join is this?

```
SELECT S.name, E.classid  
FROM Students S ?? Enrolled E  
ON S.sid=E.sid
```

S

S.name	S.sid
Jones	11111
Smith	22222
Brown	33333

S.name	E.classid
Jones	History105
Smith	French150

E

E.sid	E.classid
11111	History105
11111	DataScience194
22222	French150
44444	English10

SQL Joins

```
SELECT S.name, E.classid
FROM Students S LEFT SEMI JOIN Enrolled E
ON S.sid=E.sid
```

S

S.name	S.sid
Jones	11111
Smith	22222
Brown	33333

S.name	E.classid
Jones	History105
Smith	French150

E

E.sid	E.classid
11111	History105
11111	DataScience194
22222	French150
44444	English10

What kind of Join is this?

SELECT *

FROM **Students S ?? Enrolled E**

S

S.name	S.sid
Jones	11111
Smith	22222

E

E.sid	E.classid
11111	History105
11111	DataScience194
22222	French150

S.name	S.sid	E.sid	E.classid
Jones	11111	11111	History105
Jones	11111	11111	DataScience194
Jones	11111	22222	French150
Smith	22222	11111	History105
Smith	22222	11111	DataScience194
Smith	22222	22222	French150

SQL Joins

SELECT *

FROM **students S CROSS JOIN Enrolled E**

S

S.name	S.sid
Jones	11111
Smith	22222

E

E.sid	E.classid
11111	History105
11111	DataScience194
22222	French150

S.name	S.sid	E.sid	E.classid
Jones	11111	11111	History105
Jones	11111	11111	DataScience194
Jones	11111	22222	French150
Smith	22222	11111	History105
Smith	22222	11111	DataScience194
Smith	22222	22222	French150

What kind of Join is this?

SELECT *

FROM Students S, Enrolled E

WHERE S.sid <= E.sid

S

S.name	S.sid
Jones	11111
Smith	22222

E

E.sid	E.classid
11111	History105
11111	DataScience194
22222	French150

S.name	S.sid	E.sid	E.classid
Jones	11111	11111	History105
Jones	11111	11111	DataScience194
Jones	11111	22222	French150
Smith	22222	22222	French150

Theta Joins

```
SELECT *  
FROM Students S, Enrolled E  
WHERE S.sid <= E.sid
```

S

S.name	S.sid
Jones	11111
Smith	22222

E

E.sid	E.classid
11111	History105
11111	DataScience194
22222	French150

S.name	S.sid	E.sid	E.classid
Jones	11111	11111	History105
Jones	11111	11111	DataScience194
Jones	11111	22222	French150
Smith	22222	22222	French150

OTHER “TABLE-LIKE” DATA MODELS

Pandas/Python

- **Series:** a named, ordered dictionary
 - The keys of the dictionary are the **indexes**
 - Built on NumPy's **ndarray**
 - Values can be any Numpy data type object
- **DataFrame:** a table with named columns
 - Represented as a Dict (col_name -> series)
 - Each Series object represents a column

Operations

- map() functions
- filter (apply predicate to rows)
- sort/group by
- aggregate: sum, count, average, max, min
- Pivot or reshape
- Relational:
 - union, intersection, difference, cartesian product (CROSS JOIN)
 - select/filter, project
 - join: natural join (INNER JOIN), theta join, semi-join, etc.
 - rename

Matrices vs Databases

- Tools like Pandas give up some of the important safety features of RDBMS (e.g. ACID), but can be much faster.

Matrix multiply in SQL:

A

row	col	value
1	1	5.7
3	1	3.2
2	2	-5

B

row	col	value
2	1	12.0
3	3	5.1

```
SELECT A.row, B.col, SUM(A.value * B.value)  
FROM A JOIN B  
ON A.col = B.row  
GROUP BY A.row, B.col
```

You probably never want to do this, but the ***opposite*** direction (relational aggregate query → matrix mult.) can be very useful.

What's Wrong with Tables?



- Too limited in structure?
- Too rigid?
- Too old fashioned?

What's Wrong with (RDBMS) Tables?

- **Indices:** Typical RDBMS table storage is mostly indices
 - Can't afford this overhead for large datastores
- **Transactions:**
 - Safe state changes require journals etc., and are slow
- **Relations:**
 - Checking relations adds further overhead to updates
- **Sparse Data Support:**
 - RDBMS Tables are very wasteful when data is very sparse
 - Very sparse data is common in modern data stores
 - RDBMS tables might have dozens of columns, modern data stores might have many thousands.

RDBMS tables – row based

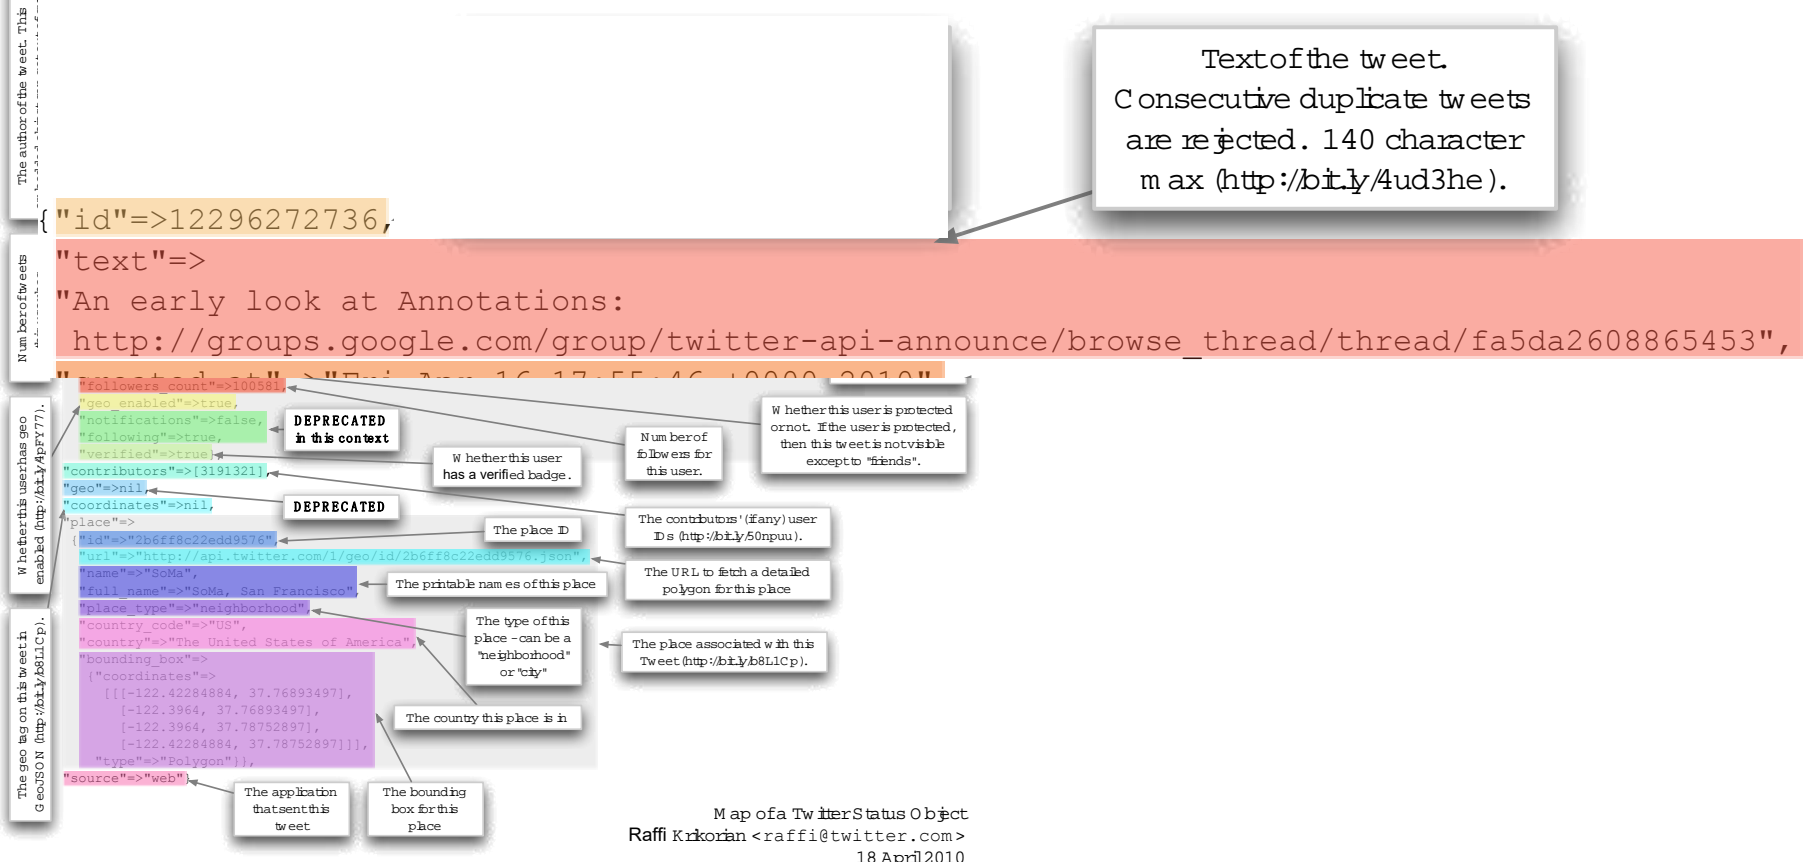
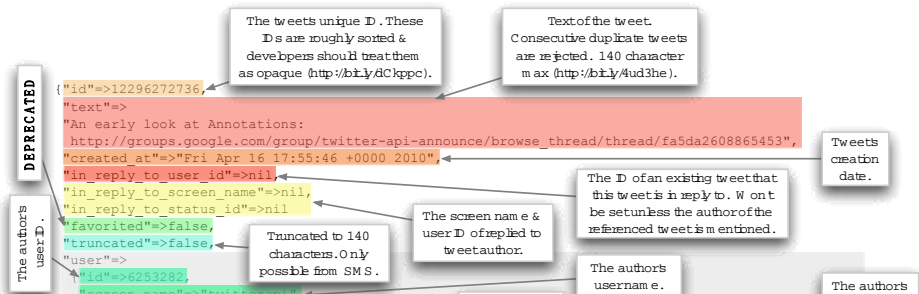
Table:

sid	name	login	age	gpa
53831	Jones	jones@cs	18	3.4
53831	Smith	smith@ee	18	3.2

Represented as:

53831	Jones	jones@cs	18	3.4
53831	Smith	smith@ee	18	3.2

Tweet JSON Format



RDBMS tables – row based

Table:

ID	name	login	loc	locid	LAT	LONG	ALT	State
52841	Jones	jones@cs	NULL	NULL	NULL	NULL	NULL	NULL
53831	Smith	smith@ee	NULL	NULL	NULL	NULL	NULL	NULL
55541	Brown	brown@ee	NULL	NULL	NULL	NULL	NULL	NULL

Represented as:

52841	Jones	jones@cs	NULL	NULL	NULL	NULL	NULL	NULL
53831	Smith	smith@ee	NULL	NULL	NULL	NULL	NULL	NULL
55541	Brown	brown@ee	NULL	NULL	NULL	NULL	NULL	NULL

Column-based store

Table:

ID	name	login	loc	locid	LAT	LONG	ALT	State
52841	Jones	jones@cs	Albany	2341	38.4	122.7	100	CA
53831	Smith	smith@ee	NULL	NULL	NULL	NULL	NULL	NULL
55541	Brown	brown@ee	NULL	NULL	NULL	NULL	NULL	NULL

Represented as column (key-value) stores:

ID	name
52841	Jones
53831	Smith
55541	Brown

ID	login
52841	jones@cs
53831	smith@ee
55541	brown@ee

ID	loc
52841	Albany

ID	locid
52841	2341

ID	LAT
52841	38.4

ID	LONG
52841	122.7

...



BEYOND TABLES

NoSQL Storage Systems



	Data Model
Cassandra	Columnfamily
CouchDB	Document
HBase	Columnfamily
MongoDB	Document
Neo4J	Graph
Redis	Collection
Riak	Document
Scalaris	Key/value
Tokyo Cabinet	Key/value
Voldemort	Key/value

Column-Family Stores (Cassandra)

A column-family groups data columns together, and is analogous to a table.

Static column family from Apache Cassandra:

row key	columns ...			
jbellis	name	email	address	state
	jonathan	jb@ds.com	123 main	TX
dhutch	name	email	address	state
	daria	dh@ds.com	45 2 nd St.	CA
egilmore	name	email		
	eric	eg@ds.com		

Dynamic Column family (Cassandra):

row key	columns ...			
jbellis	dhutch	egilmore	datastax	mzcassie
dhutch	egilmore			
egilmore	datastax	mzcassie		

CouchDB Data Model (JSON)

- “With CouchDB, no schema is enforced, so new document types with new meaning can be safely added alongside the old.”
- A CouchDB document is an object that consists of named fields. Field values may be:
 - strings, numbers, dates,
 - ordered lists, associative maps

```
"Subject": "I like Plankton"  
"Author": "Rusty"  
"PostedDate": "5/23/2006"  
"Tags": ["plankton", "baseball", "decisions"]  
"Body": "I decided today that I don't like baseball. I like plankton."
```

Prerequisites for “Schemaless” DBs

- Need **external** and **internal** representations for all data types that will be used.
- **Internal:** a dynamically-typed, object-oriented language (like Java)
- **External:** an extensible data description language: JSON or XML
- **For Performance:** Fast SerDe (Serialization and DeSerialization) so internal data structures can be efficiently pushed or extracted from disk or network.

JSON format

```
{ "firstName": "John",  
  "lastName": "Smith",  
  "isAlive": true,  
  "age": 25,  
  "height_cm": 167.6,  
  "address": {  
    "streetAddress": "21 2nd Street",  
    "city": "New York",  
    "state": "NY",  
    "postalCode": "10021-3100"  
  }  
}
```

Prerequisites for “Schemaless” DBs

- JSON includes named fields in a tree structure. Primitive types (e.g. string, number, boolean,...) are implicit.
- We can read JSON data (or XML) and automatically create internal representations for complex data.
- Using the **field names** and **object structure**, we can query these objects once loaded.

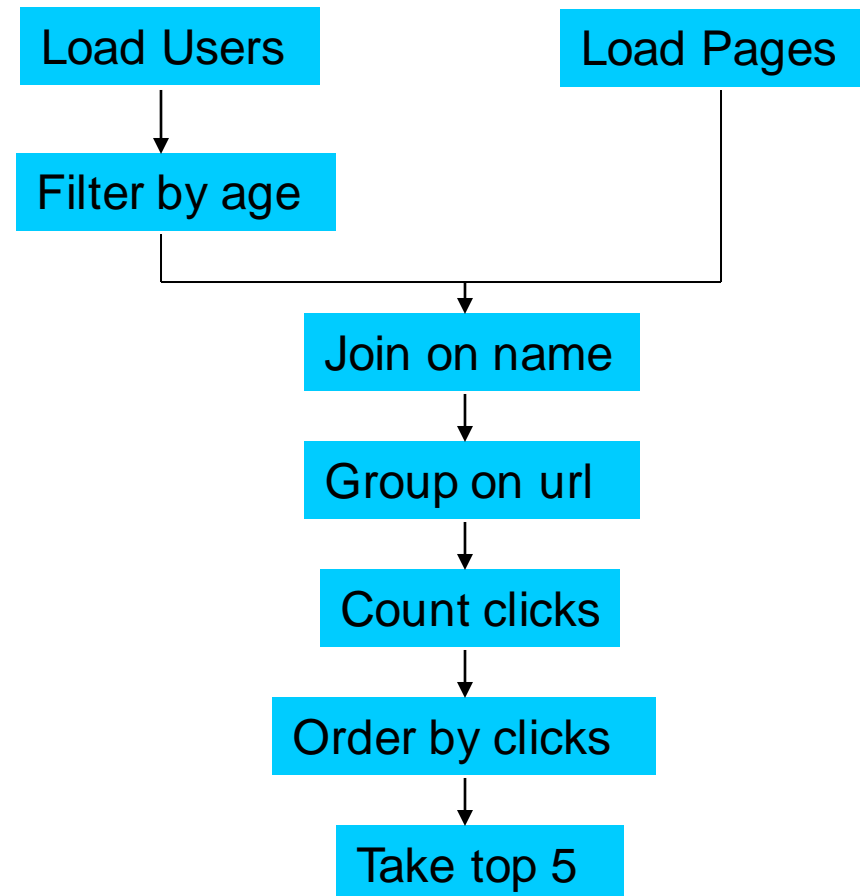
Pig

- Started at Yahoo! Research
- Runs about 50% of Yahoo!'s jobs
- Features:
 - Expresses sequences of MapReduce jobs
 - Data model: nested “bags” of items
 - Schema is optional
 - Provides relational (SQL) operators (JOIN, GROUP BY, etc)
 - Easy to plug in Java functions



An Example Problem

Suppose you have user data in one file, website data in another, and you need to find the top 5 most visited pages by users aged 18-25.



In MapReduce

```
import java.io.IOException;
import java.util.ArrayList;
import java.util.Iterator;
import java.util.List;

import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.io.WritableComparable;
import org.apache.hadoop.mapred.FileInputFormat;
import org.apache.hadoop.mapred.FileOutputFormat;
import org.apache.hadoop.mapred.JobConf;
import org.apache.hadoop.mapred.Mapper;
import org.apache.hadoop.mapred.Reducer;
import org.apache.hadoop.mapred.Reporter;
import org.apache.hadoop.mapred.SequenceFileInputFormat;
import org.apache.hadoop.mapred.SequenceFileOutputFormat;
import org.apache.hadoop.mapred.TextInputFormat;
import org.apache.hadoop.mapred.JobControl;
import org.apache.hadoop.mapred.lib.IdentityMapper;

public class MRExample {
    public static class LoadPages extends MapReduceBase
        implements Mapper<LongWritable, Text, Text, Text> {

        public void map(LongWritable k, Text val,
            OutputCollector<Text, Text> oc,
            Reporter reporter) throws IOException {
            // Pull the key out
            String line = val.toString();
            int firstComma = line.indexOf(',');
            String key = line.substring(0, firstComma);
            String value = line.substring(firstComma + 1);
            Text outKey = new Text(key);
            // Prepend an index to the value so we know which file
            // it came from.
            Text outVal = new Text("1 " + value);
            oc.collect(outKey, outVal);
        }

        public static class LoadAndFilterUsers extends MapReduceBase
            implements Mapper<LongWritable, Text, Text, Text> {

        public void map(LongWritable k, Text val,
            OutputCollector<Text, Text> oc,
            Reporter reporter) throws IOException {
            // Pull the key out
            String line = val.toString();
            int firstComma = line.indexOf(',');
            String value = line.substring(firstComma + 1);
            int age = Integer.parseInt(value);
            if (age < 18 || age > 25) return;
            String key = line.substring(0, firstComma);
            Text outKey = new Text(key);
            // Prepend an index to the value so we know which file
            // it came from.
            Text outVal = new Text("2 " + value);
            oc.collect(outKey, outVal);
        }

        public static class Join extends MapReduceBase
            implements Reducer<Text, Text, Text, Text> {

        public void reduce(Text key,
            Iterator<Text> iter,
            OutputCollector<Text, Text> oc,
            Reporter reporter) throws IOException {
            // For each value, figure out which file it's from and
            // accordingly.
            List<String> first = new ArrayList<String>();
            List<String> second = new ArrayList<String>();

            while (iter.hasNext()) {
                Text t = iter.next();
                String value = t.toString();
                if (value.charAt(0) == '1')
                    first.add(value.substring(1));
                else second.add(value.substring(1));
            }

            reporter.setStatus("OK");
        }

        // Do the cross product and collect the values
        for (String s1 : first) {
            for (String s2 : second) {
                String outval = key + "," + s1 + "," + s2;
                oc.collect(null, new Text(outval));
                reporter.setStatus("OK");
            }
        }
    }

    public static class LoadJoined extends MapReduceBase
        implements Mapper<Text, Text, Text, LongWritable> {

        public void map(
            Text k,
            Text val,
            OutputCollector<Text, LongWritable> oc,
            Reporter reporter) throws IOException {
            // Find the url
            String line = val.toString();
            int firstComma = line.indexOf(',');
            int secondComma = line.indexOf(',', firstComma);
            String key = line.substring(firstComma, secondComma);
            // drop the rest of the record, I don't need it anymore,
            // just pass a 1 for the combiner/reducer to sum instead.
            Text outKey = new Text(key);
            Text outVal = new LongWritable(1L);
            oc.collect(outKey, outVal);
        }

        public static class ReduceUrls extends MapReduceBase
            implements Reducer<Text, LongWritable, WritableComparable,
                Writable> {

        public void reduce(
            Text key,
            Iterator<LongWritable> iter,
            OutputCollector<WritableComparable, Writable> oc,
            Reporter reporter) throws IOException {
            // Add up all the values we see
            long sum = 0;
            while (iter.hasNext()) {
                sum += iter.next().get();
                reporter.setStatus("OK");
            }

            oc.collect(key, new LongWritable(sum));
        }

        public static class LoadClicks extends MapReduceBase
            implements Mapper<WritableComparable, Writable, LongWritable,
                Text> {

        public void map(
            WritableComparable key,
            Writable val,
            OutputCollector<LongWritable, Text> oc,
            Reporter reporter) throws IOException {
                oc.collect((LongWritable)val, (Text)key);
            }

        public static class LimitClicks extends MapReduceBase
            implements Reducer<LongWritable, Text, LongWritable, Text> {

        int count = 0;

        public void reduce(
            LongWritable key,
            Iterator<Text> iter,
            OutputCollector<LongWritable, Text> oc,
            Reporter reporter) throws IOException {
                // Only output the first 100 records
                while (count < 100 && iter.hasNext()) {
                    oc.collect(key, iter.next());
                    count++;
                }
            }

        public static void main(String[] args) throws IOException {
            JobConf lp = new JobConf(MRExample.class);
            lp.setJobName("Load Pages");
            lp.setInputFormat(TextInputFormat.class);
            lp.setOutputKeyClass(Text.class);
            lp.setOutputValueClass(Text.class);
            lp.setMapperClass(LoadPages.class);
            FileInputFormat.addInputPath(lp, new
                Path("/user/gates/pages"));
            FileOutputFormat.setOutputPath(lp,
                new Path("/user/gates/tmp/indexed_pages"));
            lp.setNumReduceTasks(0);
            Job loadPages = new Job(lp);

            JobConf ifu = new JobConf(MRExample.class);
            ifu.setJobName("Load and Filter Users");
            ifu.setInputFormat(TextInputFormat.class);
            ifu.setOutputKeyClass(Text.class);
            ifu.setOutputValueClass(Text.class);
            ifu.setMapperClass(LoadAndFilterUsers.class);
            FileInputFormat.addInputPath(ifu, new
                Path("/user/gates/users"));
            FileOutputFormat.setOutputPath(ifu,
                new Path("/user/gates/tmp/filtered_users"));
            ifu.setNumReduceTasks(0);
            Job loadUsers = new Job(ifu);

            JobConf join = new JobConf(MRExample.class);
            join.setJobName("Join Users and Pages");
            join.setInputFormat(KeyValueTextInputFormat.class);
            join.setOutputKeyClass(Text.class);
            join.setOutputValueClass(Text.class);
            join.setMapperClass(IdentityMapper.class);
            join.setReducerClass(Join.class);
            FileInputFormat.addInputPath(join, new
                Path("/user/gates/tmp/indexed_pages"));
            FileInputFormat.addInputPath(join, new
                Path("/user/gates/tmp/filtered_users"));
            FileOutputFormat.setOutputPath(join, new
                Path("/user/gates/tmp/joined"));
            join.setNumReduceTasks(50);
            Job joinJob = new Job(join);
            joinJob.addDependingJob(loadPages);
            joinJob.addDependingJob(loadUsers);

            JobConf group = new JobConf(MRExample.class);
            group.setJobName("Group URLs");
            group.setInputFormat(KeyValueTextInputFormat.class);
            group.setOutputKeyClass(Text.class);
            group.setOutputValueClass(LongWritable.class);
            group.setOutputFormat(SequenceFileOutputFormat.class);
            group.setMapperClass(LoadJoined.class);
            group.setCombinerClass(ReduceUrls.class);
            group.setReducerClass(ReduceUrls.class);
            FileInputFormat.addInputPath(group, new
                Path("/user/gates/tmp/joined"));
            FileOutputFormat.setOutputPath(group, new
                Path("/user/gates/tmp/grouped"));
            group.setNumReduceTasks(50);
            Job groupJob = new Job(group);
            groupJob.addDependingJob(joinJob);

            JobConf top100 = new JobConf(MRExample.class);
            top100.setJobName("Top 100 sites");
            top100.setInputFormat(SequenceFileInputFormat.class);
            top100.setOutputKeyClass(LongWritable.class);
            top100.setOutputValueClass(Text.class);
            top100.setOutputFormat(SequenceFileOutputFormat.class);
            top100.setMapperClass(LoadClicks.class);
            top100.setCombinerClass(LimitClicks.class);
            top100.setReducerClass(LimitClicks.class);
            FileInputFormat.addInputPath(top100, new
                Path("/user/gates/tmp/grouped"));
            FileOutputFormat.setOutputPath(top100, new
                Path("/user/gates/top100sitesforusers18to25"));
            top100.setNumReduceTasks(1);
            Job limit = new Job(top100);
            limit.addDependingJob(groupJob);

            JobControl jc = new JobControl("Find top 100 sites for users
                18 to 25");
            jc.addJob(loadPages);
            jc.addJob(loadUsers);
            jc.addJob(joinJob);
            jc.addJob(groupJob);
            jc.addJob(limit);
            jc.run();
        }
    }
}
```

In Pig Latin

```
Users      = load 'users' as (name, age);
Filtered   = filter Users by
              age >= 18 and age <= 25;
Pages      = load 'pages' as (user, url);
Joined     = join Filtered by name, Pages by user;
Grouped    = group Joined by url;
Summed     = foreach Grouped generate group,
              count(Joined) as clicks;
Sorted     = order Summed by clicks desc;
Top5       = limit Sorted 5;

store Top5 into 'top5sites';
```

Hive

- Developed at Facebook
- Relational database built on Hadoop
 - Maintains table schemas
 - SQL-like query language (which can also call Hadoop Streaming scripts)
 - Supports table partitioning, complex data types, sampling, some query optimization
- Used for most Facebook jobs
 - Less than 1% of daily jobs at Facebook use MapReduce directly!!! (SQL – or PIG – wins!)
 - Note: Google also has several SQL-like systems in use.



DATASPACE – WHAT ARE THEY?

Dataspaces

Inclusive

Deal with **all** the data of interest – in whatever form

Co-existence not Integration

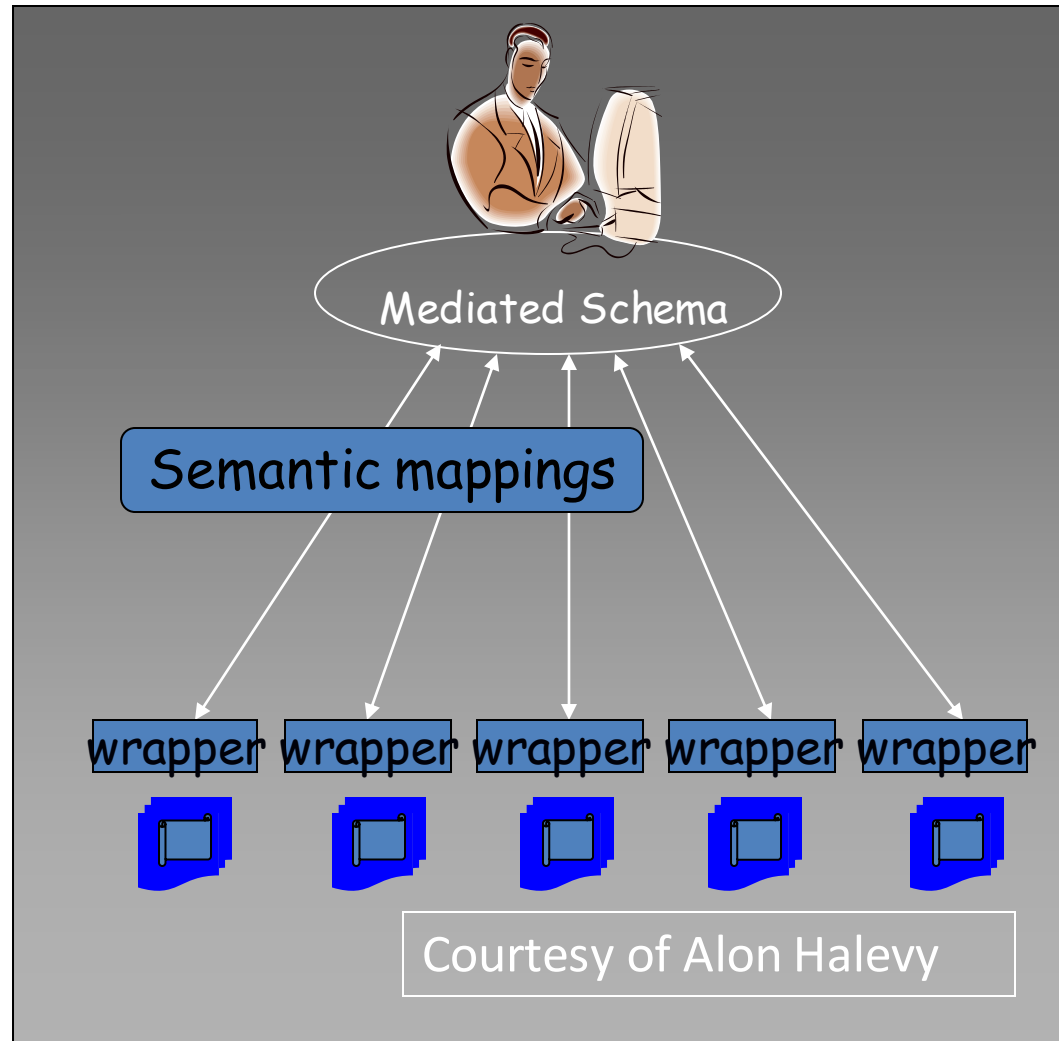
No integrated schema, no single warehouse, no ownership required

Pay-as-you-go

- Keyword search is bare minimum.
- More function and increased consistency as you add work.

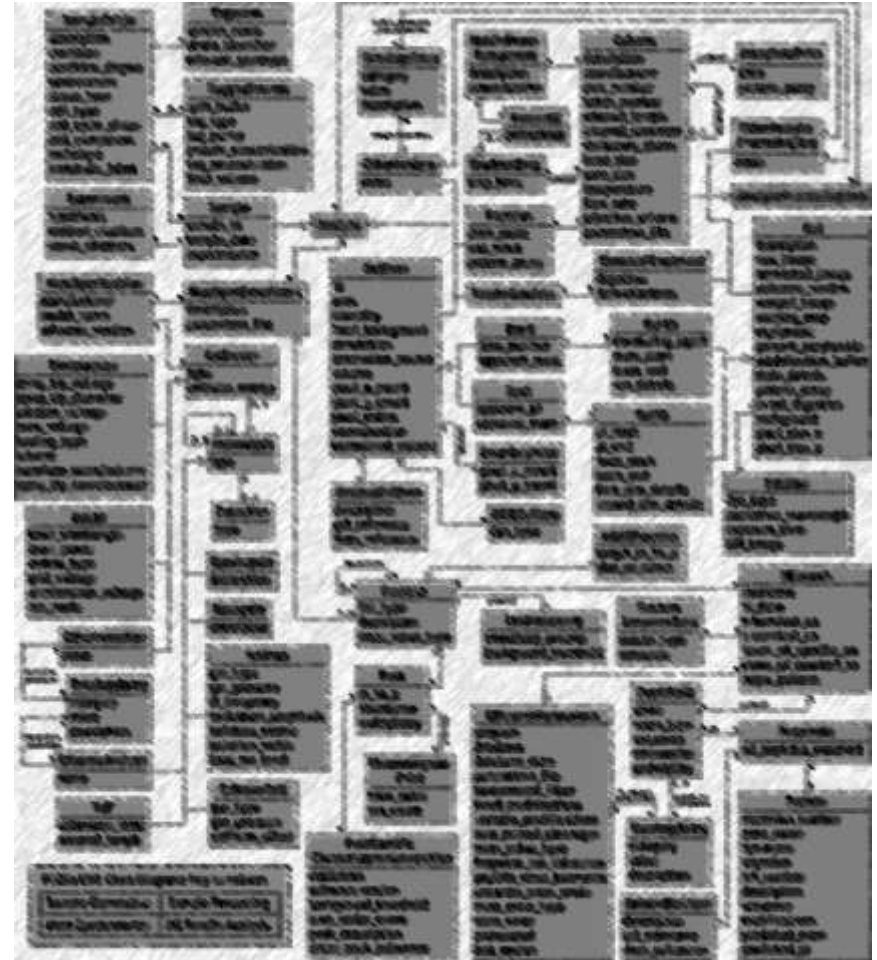
Compare to Data Integration

A quintessential
schema-first
approach.

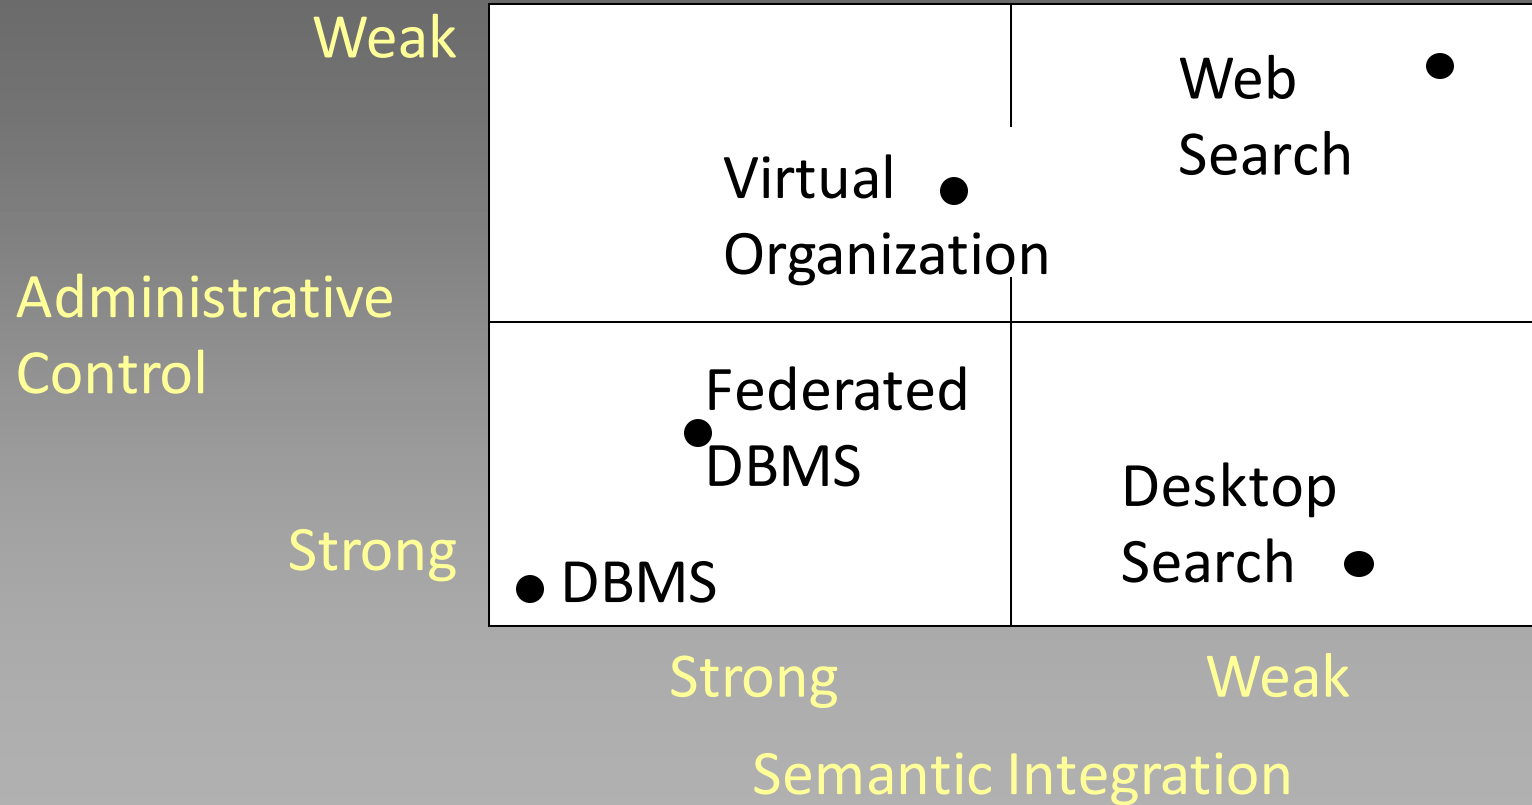


Whither Structured Data?

- Conventional Wisdom:
only 20% of data is structured.
- Decreasing due to:
 - Consumer applications
 - Enterprise search
 - Media applications

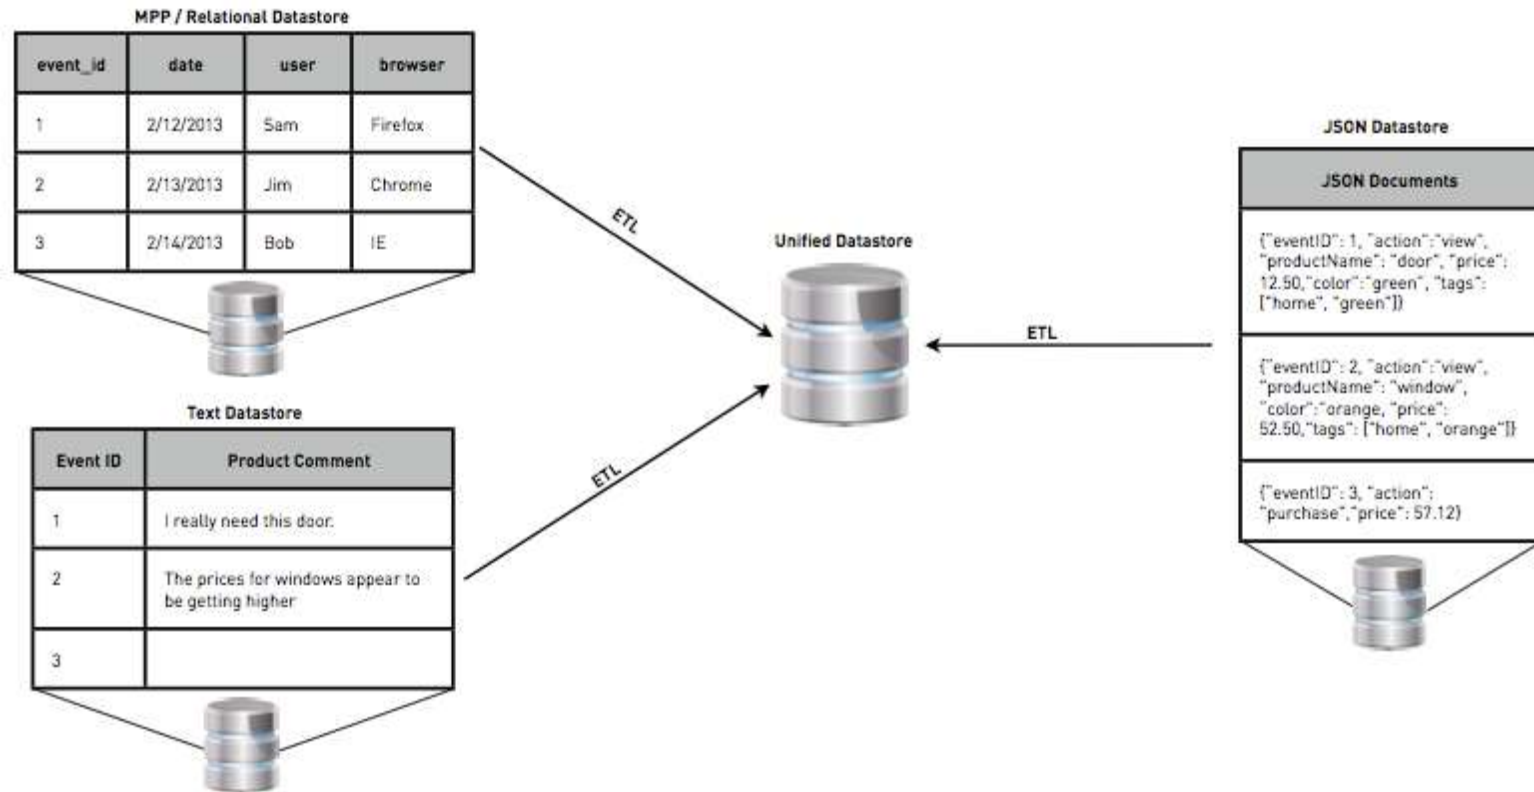


An Alternative View



A Recent Example

- Hadapt's “Schemaless SQL”



A Recent Example

- Hadapt's "Schemaless SQL"

actions

date	user	browser	json	product_comment
2/12/2013	Sam	Firefox	{"eventID": 1, "action": "view", "productName": "door", "price": 12.50, "color": "green", "tags": ["home", "green"]}	I really need this door.
2/13/2013	Jim	Chrome	{"eventID": 2, "action": "view", "productName": "window", "color": "orange", "price": 52.50, "tags": ["home", "orange"]}	The prices for windows appear to be getting higher.
2/14/2013	Bob	IE	{"eventID": 3, "action": "purchase", "price": 57.12}	

actions

date	user	browser	product name	price	tags	color	action	product comment
2/12/2013	Sam	Firefox	door	12.50	home green	green	view	I really need this door.
2/13/2013	Jim	Chrome	window	52.50	home orange	orange	view	The prices for windows appear to be going higher.
2/14/2013	Bob	IE		57.12			purchase	

“Schemaless SQL”

- Schema Evolution – adding a column

date	user	browser	json	product_comment
2/15/2013	Jim	Firefox	[{"eventID": 4, "action": "view", "productName": "door", "price": 12.50, "color": "green", "tags": ["home", "green"], "size": "standard"}]	I need a standard sized door.

```
SELECT user, product_name, size
FROM actions
WHERE action='view'
AND date > 2/1/2013
AND matches('product_comment:higher')
```

actions									
date	user	browser	product_name	price	tags	color	action	product comment	size
2/12/2013	Sam	Firefox	door	12.50	home green	green	view	I really need this door.	
2/13/2013	Jim	Chrome	window	52.50	home orange	orange	view	The prices for windows appear to be going higher.	
2/14/2013	Bob	IE		57.12			purchase		
2/15/2013	Jim	Firefox	door	12.50	home green	green	view	I need a standard sized door.	standard

Not “IF” But “WHEN”?

- “Schema on Write”
 - Traditional Approach
- “Schema on Read”
 - Data is simply copied to the file store, no transformation is needed.
 - A SerDe (Serializer/Deserializer) is applied during read time to extract the required columns (late binding)
 - New data can start flowing anytime and will appear retroactively once the SerDe is updated to parse it.

• Read is Fast

• Standards/Governance



• Load is Fast

• Flexibility/Agility

Summary

- Data Models, Tables, Structure, etc.
 - SQL
 - NoSQL
 - Schema on Read vs. Schema on Write